

HW1 Peer Assessment

Part A. ANOVA

Additional Material: ANOVA tutorial

<https://data-science-plus.com/one-way-anova-in-r/>

Jet lag is a common problem for people traveling across multiple time zones, but people can gradually adjust to the new time zone since the exposure of the shifted light schedule to their eyes can reset the internal circadian rhythm in a process called "phase shift". Campbell and Murphy (1998) in a highly controversial study reported that the human circadian clock can also be reset by only exposing the back of the knee to light, with some hairing this as a major discovery and others challenging aspects of the experimental design. The table below is taken from a later experiment by Krings and Conder (2002) that re-examined the phenomenon. The new experiment measured circadian rhythm through the daily cycle of melatonin production in 22 subjects randomly assigned to one of three light treatments. Subjects were woken from sleep and for three hours were exposed to bright lights applied to the eyes only, to the knees only or to neither (control group). The effects of treatment to the circadian rhythm were measured two days later by the magnitude of phase shift (measured in hours) in each subject's daily cycle of melatonin production. A negative measurement indicates a delay in melatonin production, a predicted effect of light treatment, while a positive number indicates an advance.

Raw data of phase shift, in hours, for the circadian rhythm experiment

Treatment	Phase Shift (hr)
Control	0.53, 0.36, 0.20, -0.37, -0.60, -0.64, -0.68, -1.27
Eyes	0.73, 0.31, 0.03, 0.29, -0.56, -0.96, -1.63
Knees	-0.76, -0.86, -1.35, -1.48, -1.52, -2.04, -2.83

Creating data

```
treatment <- c("Control","Control","Control","Control","Control","Control","Control","Control",
              "Knees","Knees","Knees","Knees","Knees","Knees","Knees",
              "eyes","eyes","eyes","eyes","eyes","eyes","eyes","eyes")
shift <- c(-0.53, 0.36, 0.20, -0.37, -0.60, -0.64, -0.68, -1.27,
          0.73, 0.31, 0.03, -0.29, -0.56, -0.96, -1.63,
          -0.76, -0.86, -1.35, -1.48, -1.52, -2.04, -2.83)
df <- data.frame(treatment, shift)
print(df)

## treatment shift
## 1 Control 0.53
## 2 Control 0.36
## 3 Control 0.20
## 4 Control -0.37
## 5 Control -0.68
## 6 Control -0.64
## 7 Control -0.68
## 8 Control -1.27
## 9 Knees -0.72
## 10 Knees 0.31
## 11 Knees 0.83
## 12 Knees -0.29
## 13 Knees -0.56
## 14 Knees -0.96
## 15 Knees -1.61
## 16 Eyes -0.78
## 17 Eyes -0.86
## 18 Eyes -1.35
## 19 Eyes -1.48
## 20 Eyes -1.52
## 21 Eyes -2.84
## 22 Eyes -1.83
```

Question A1 - 3 pts

Consider the following incomplete R output.

Source	Df	Sum of Squares	Mean Squares	F-statistics	p-value
Treatments	?	?	3.6122	?	0.004
Error	?	9.415	?		
TOTAL	?	?			

Fill in the missing values in the analysis of the variance table. Note: Missing values can be calculated using the corresponding formulas provided in the lectures, or you can build the data frame in R and generate the ANOVA table using the aov() function. Either approach will be accepted.

```
model <- aov(shift ~ treatment, data = df)
summary(model)

##           Df Sum Sq Mean Sq F value    Pr(>F)
## treatment    2  7.224   3.612   7.289 0.00447 **
## Residuals   19  9.415   0.496
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

model.tables(model3, types="means")

## Tables of means
## Grand mean
##
## -0.727273
##
## treatment
## Control Eyes Knees
## -0.3887 -1.551 -0.3357
## rep    8.0880  7.088  7.0880

## Source      Df      Sum of Squares      Mean Squares      F-statistics      p-value
## Treatments    2          7.224          3.6122          7.289          0.004
## Error         19          9.415          0.496
## TOTAL         21         16.639
```

Question A2 - 3 pts

Use μ_1 , μ_2 , and μ_3 as notation for the three mean parameters and define these parameters clearly based on the context of the topic above (i.e. explain what μ_1 , μ_2 , and μ_3 mean in words in the context of this problem). Find the estimates of these parameters.

```
model.tables(model3, types="means")

## Tables of means
## Grand mean
##
## -0.727273
##
## treatment
## Control Eyes Knees
## -0.3887 -1.551 -0.3357
## rep    8.0880  7.088  7.0880

mu1: Average phase shift (measured in hours) in each subject's daily cycle of melatonin production for subjects in the control group (no light) Value: -0.3887 hours
mu2: Average phase shift (measured in hours) in each subject's daily cycle of melatonin production for subjects in the eyes group Value: -1.551 hours
mu3: Average phase shift (measured in hours) in each subject's daily cycle of melatonin production for subjects in the knees group Value: -0.3357 hours
```

Question A3 - 5 pts

Use the ANOVA table in Question A1 to answer the following questions:

- 1 pts Write the null hypothesis of the ANOVA F -test, H_0 . H_0 : All three groups (2 treatments + control) will have the same mean effect on phase shift
- 1 pts Write the alternative hypothesis of the ANOVA F -test, H_1 . H_1 : At least 1 of the three groups (2 treatments + control) will have a different mean effect on phase shift
- 1 pts Fill in the blanks for the degrees of freedom of the ANOVA F -test statistic:
 $DF(2, 19)$
- 1 pts What is the p-value of the ANOVA F -test? $p = .00447$
- 1 pts According to the results of the ANOVA F -test, does light treatment affect phase shift? Use an α -level of 0.05. Yes

Part B. Simple Linear Regression

We are going to use regression analysis to estimate the performance of CPUs based on the maximum number of channels in the CPU. This data set comes from the UCI Machine Learning Repository.

The data file includes the following columns:

- vendor: vendor of the CPU
- chmax: maximum channels used in the CPU
- performance: published relative performance of the CPU

The data is in the file "machine.csv". To read the data in R, save the file in your working directory (make sure you have changed the directory if different from the R working directory) and read the data using the R function read.csv().

```
# Read in the data
data <- read.csv("machine.csv", head = TRUE, sep = ";")
# Show the first few rows of data
head(data, 3)

## vendor      chmax      performance
## <chr>      <dbl>      <dbl>
##
## 1  amd64      128      198
## 2  amd64      32      269
## 3  amd64      32      220
## #> rows
```

Question B1: Exploratory Data Analysis - 9 pts

- 3 pts Use a scatter plot to describe the relationship between CPU performance and the maximum number of channels. Describe the general trend (direction and form). Include plots and R-code used.

```
attach(data)
plot(chmax, performance)

## There seems to be a positive relationship between performance and max # of channels, though there are some outliers and/or leverage points that have chmax much higher than the bulk of the group. It also appears the data is a "cone" shape so the data is not linear, meaning that chmax alone does not explain all of the variability in performance.

## b. 3 pts What is the value of the correlation coefficient between performance and chmax? Please interpret the strength of the correlation based on the correlation coefficient.

## Your code here...
model <- lm(performance ~ chmax)
summary(model)

## Call:
## lm(formula = performance ~ chmax)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -486.47  -42.20  -22.20   20.31  867.35
##
## Coefficients:
## (Intercept)  37.2252  18.8587  3.428 0.800733 ***
## chmax        3.7441   0.3423 10.938 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 128.3 on 287 degrees of freedom
## Multiple R-squared:  0.3663, Adjusted R-squared:  0.3632
## F-statistic: 119.6 on 1 and 287 DF, p-value: < 2.2e-16

## Your code here...
cor(chmax, performance)

## [1] 0.6852893
```

The correlation coefficient is 0.6852. This is a fairly strong positive correlation.

- 2 pts Based on this exploratory analysis, would you recommend a simple linear regression model for the relationship? I would say no. The coefficient of determination (r -squared) is .3663, indicating that only one third of the variability in performance can be explained by variability in chmax, which is not that good if it will be used as the sole predictor.
- 1 pts Based on the analysis above, would you pursue a transformation of the data? Do not transform the data. Yes, the funnel shape and the clustering of points at the lower end of chmax suggests to me that a logarithmic transformation may compress the data and lead to a more linear relationship!

Question B2: Fitting the Simple Linear Regression Model - 11 pts

Fit a linear regression model, named model1, to evaluate the relationship between performance and the maximum number of channels. Do not transform the data. The function you should use is fit.R:

```
# Your code here...
model1 <- lm(performance ~ chmax, data)
summary(model1)

## Call:
## lm(formula = performance ~ chmax, data = data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -486.47  -42.20  -22.20   20.31  867.35
##
## Coefficients:
## (Intercept)  37.2252  18.8587  3.428 0.800733 ***
## chmax        3.7441   0.3423 10.938 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 128.3 on 287 degrees of freedom
## Multiple R-squared:  0.3663, Adjusted R-squared:  0.3632
## F-statistic: 119.6 on 1 and 287 DF, p-value: < 2.2e-16

## a. 3 pts What are the model parameters and what are their estimates?
## The model parameters are chmax and the intercept. The estimate of chmax is 3.7441 and the estimate of the intercept is 37.2252

## b. 2 pts Write down the estimated simple linear regression equation.  $y = 3.7441x_{chmax} + 37.2252 + e$ 

## c. 2 pts Interpret the estimated value of the  $\beta_1$  parameter in the context of the problem.  $\beta_1$ , which is chmax in this case, is 3.7441. This suggests that for every additional maximum channel in the CPU, the performance metric will improve by 3.7441.

## d. 2 pts Find a 95% confidence interval for the  $\beta_1$  parameter. Is  $\beta_1$  statistically significant at this level?

## confint(model1, level=.95)

##           2.5 %      97.5 %
## (Intercept) 15.81792 58.623848
## chmax       3.69251  4.438905

## The 95% confidence interval for  $\beta_1$  is (3.69251, 4.438905). That makes  $\beta_1$  statistically significant (non zero) at that level.

## e. 2 pts Is  $\beta_1$  statistically significantly positive at an  $\alpha$ -level of 0.01? What is the approximate p-value of this test?

## tvalue=10.938
## p=(tvalue,287)

## [1] 0

## confInt(model1, level=.95)
```

The p-value of the one sided test is so small that R interprets it as 0, so yes the β_1 is statistically significantly positive at an α -level of 0.01

Question B3: Checking the Assumptions of the Model - 8 pts

Create and interpret the following graphs with respect to the assumptions of the linear regression model. In other words, comment on whether there are any apparent departures from the assumptions of the linear regression model. Make sure that you state the model assumptions and assess each one. Each graph may be used to assess one or more model assumptions.

- 2 pts Scatterplot of the data with chmax on the x-axis and performance on the y-axis

```
plot(chmax, performance)

## Fitted vs Residuals

## Model Assumption(s) checks: Linearity/mean 0 assumption
## Interpretation: It appears generally linear, but it is imperfect.

## b. 3 pts Residual plot - a plot of the residuals,  $\epsilon_i$ , versus the fitted values,  $\hat{y}_i$ 

## plot(fitted(model1), resid(model1), main="Fitted vs Residuals", xlab="Fitted Values")

## Fitted vs Residuals

## Model Assumption(s) checks: Constant variance assumption/ Independence assumption (not really, we can only check for uncorrelated errors)
## Interpretation: There is no clear shape to the distribution (such as a curve) that suggests there is a non-linear relationship, but we see a megaphone effect on the residuals, which means the constant variance assumption does not hold. There are no clusters of residuals, which suggests the errors are not correlated (related to but not proof of independence assumption)
```

- 3 pts Histogram and q-q plot of the residuals

```
qqnorm(resid(model1), main="QQ-Plot of Residuals")
qqline(resid(model1))

## QQ-Plot of Residuals

## Model Assumption(s) checks:
## Interpretation: The curvature at the ends suggests non-normality. The histogram suggests a non-normal distribution that is tall at the center with long, thin tails.

## hist(resid(model1), breaks=28)

## Histogram of resid(model1)

## Model Assumption(s) checks:
## Interpretation: The curvature at the ends suggests non-normality. The histogram suggests a non-normal distribution that is tall at the center with long, thin tails.
```

Question B4: Improving the Fit - 10 pts

- 2 pts Use a Box-Cox transformation (boxcox() in car() package or boxcox() in MASS() package) to find the optimal λ value rounded to the nearest half integer. What transformation of the response, if any, does it suggest to perform? The optimal lambda, rounded to the nearest half integer, seems to round to 0. This suggests we use the normal logarithmic transformation according to the lecture.

```
require(car)

## Loading required package: car

## Loading required package: carData

boxcox(model1, lambdaSeq=c(-1,1,.5))

## Profile Log-likelihood

## b. 2 pts Create a linear regression model, named model2, that uses the log transformed performance as the response, and the log transformed chmax as the predictor. Note: The variable chmax has a couple of zero values which will cause problems when taking the natural log. Please add one to the predictor before taking the natural log of it

## lchmax <- log(chmax+1)
## lperf <- log(performance)
## model2 <- lm(lperf ~ lchmax, data)
## summary(model2)

## Call:
## lm(formula = lperf ~ lchmax, data = data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.22543 -0.59429  0.01865  0.59287  1.85995
##
## Coefficients:
## (Intercept)  2.47655  0.14552 17.5 <2e-16 ***
## lchmax       0.64819  16.64891 12.0 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.807 on 287 degrees of freedom
## Multiple R-squared:  0.4183, Adjusted R-squared:  0.4074
## F-statistic: 144 on 1 and 287 DF, p-value: < 2.2e-16

## c. 2 pts Compare the R-squared values of model1 and model2. Did the transformation improve the explanatory power of the model? Model 1: Multiple R-squared: 0.3663, Adjusted R-squared: 0.3632 Model 2: Multiple R-squared: 0.4103, Adjusted R-squared: 0.4074 Yes, it did improve the explanatory power somewhat.

## d. 4 pts Similar to Question B3, interpret all model assumptions of model2. A model is considered a good fit if all assumptions hold. Based on your interpretation of the model assumptions, does model2 a good fit?

## plot(lchmax, lperf)
```

```
plot(fitted(model2), resid(model2), main="Fitted vs Residuals", xlab="Fitted Values")

## Fitted vs Residuals

## qqnorm(resid(model2), main="QQ-Plot of Residuals")
## qqline(resid(model2))

## QQ-Plot of Residuals

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```

```
hist(resid(model2), breaks=28)

## Histogram of resid(model2)

## The scatterplot of log(performance)
```